

ASSIGNMENT 9.2

Problem1

create hbase table named clicks with column family hits such that it should be able to store last 5 values of qualifier inside hits family

```
hbase(main):001:0> create 'clicks','IP','hits'
```

```
SLF4J: Class path contains multiple SLF4J bindings.
```

```
SLF4J: Found binding in
```

```
[jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
```

```
SLF4J: Found binding in
```

```
[jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
```

```
SLF4J: See http://www.slf4j.org/codes.html#multiple\_bindings for an explanation.
```

```
2017-11-02 05:02:53,649 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
0 row(s) in 5.4230 seconds
```

=> Hbase::Table – clicks

```
describe 'clicks'
```

```
Table clicks is ENABLED
```

```
clicks
```

```
COLUMN FAMILIES DESCRIPTION
```

```
{NAME => 'IP', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMP
```

```
RESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'hits', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', CO
```

```
MPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
2 row(s) in 1.7510 seconds
```

```
hbase(main):003:0> put 'clicks','01.02.03.198','hits:one','100'
```

```
0 row(s) in 0.8890 seconds
```

```
hbase(main):005:0> put 'clicks','01.02.03.198','hits:two','200'
```

0 row(s) in 0.0680 seconds

```
hbase(main):006:0> put 'clicks','01.02.03.198','hits:three','300'
```

0 row(s) in 0.0710 seconds

```
hbase(main):007:0> put 'clicks','01.02.03.198','hits:four','400'
```

0 row(s) in 0.0470 seconds

```
hbase(main):008:0> put 'clicks','01.02.03.198','hits:five','500'
```

0 row(s) in 0.0220 seconds

```
hbase(main):009:0> scan 'clicks'
```

ROW	COLUMN+CELL
01.02.03.198	column=hits:five, timestamp=1509579967663, value=500
01.02.03.198	column=hits:four, timestamp=1509579957383, value=400
01.02.03.198	column=hits:one, timestamp=1509579820603, value=100
01.02.03.198	column=hits:three, timestamp=1509579944904, value=300
01.02.03.198	column=hits:two, timestamp=1509579918413, value=200

1 row(s) in 0.3110 seconds

Problem 2 updating current values

```
hbase(main):010:0> put 'clicks','01.02.03.198','hits:five','1000'
```

0 row(s) in 0.0420 seconds

```
hbase(main):011:0> put 'clicks','01.02.03.198','hits:four','2000'
```

0 row(s) in 0.0700 seconds

```
hbase(main):012:0> scan 'clicks'
```

ROW	COLUMN+CELL
01.02.03.198	column=hits:five, timestamp=1509580239017, value=1000
01.02.03.198	column=hits:four, timestamp=1509580251272, value=2000
01.02.03.198	column=hits:one, timestamp=1509579820603, value=100
01.02.03.198	column=hits:three, timestamp=1509579944904, value=300
01.02.03.198	column=hits:two, timestamp=1509579918413, value=200

1 row(s) in 0.2330 seconds

```
hbase(main):013:0> alter 'clicks',NAME=>'hits',VERSIONS=>5
```

Updating all regions with the new schema...

0/1 regions updated.

0/1 regions updated.

1/1 regions updated.

Done.

0 row(s) in 4.3680 seconds

hbase(main):014:0> put 'clicks','10.101.223.445','hits:two','220'

0 row(s) in 0.0540 seconds

hbase(main):015:0> put 'clicks','10.20.30.198','hits:five','200'

0 row(s) in 0.0590 seconds

hbase(main):016:0> put 'clicks','01.02.03','hits:five','200'

0 row(s) in 0.0890 seconds

hbase(main):017:0> put 'clicks','01.02.03','hits:five','5000'

0 row(s) in 0.0210 seconds

hbase(main):018:0> put 'clicks','01.02.03','hits:five','650'

0 row(s) in 0.0490 seconds

hbase(main):019:0> scan clicks

NameError: undefined local variable or method `clicks' for #<Object:0x750ed637>

hbase(main):020:0> scan 'clicks'

ROW	COLUMN+CELL
-----	-------------

01.02.03	column=hits:five, timestamp=1509582111964, value=650
----------	--

01.02.03.198	column=hits:five, timestamp=1509580239017, value=1000
--------------	---

01.02.03.198	column=hits:four, timestamp=1509580251272, value=2000
--------------	---

01.02.03.198	column=hits:one, timestamp=1509579820603, value=100
--------------	---

01.02.03.198	column=hits:three, timestamp=1509579944904, value=300
--------------	---

01.02.03.198	column=hits:two, timestamp=1509579918413, value=200
--------------	---

10.101.223.445	column=hits:two, timestamp=1509582010376, value=220
----------------	---

10.20.30.198	column=hits:five, timestamp=1509582048809, value=200
--------------	--

4 row(s) in 0.3150 seconds

hbase(main):021:0> scan 'clicks',{COLUMN=>'hits:five',VERSIONS=>5}

ROW	COLUMN+CELL
-----	-------------

01.02.03	column=hits:five, timestamp=1509582111964, value=650
----------	--

01.02.03	column=hits:five, timestamp=1509582100596, value=5000
01.02.03	column=hits:five, timestamp=1509582089326, value=200
01.02.03.198	column=hits:five, timestamp=1509580239017, value=1000
10.20.30.198	column=hits:five, timestamp=1509582048809, value=200

3 row(s) in 0.0860 seconds